SUBQ

Understanding AI

Capabilities, Limitations & Opportunities for Media

A Briefing for Media Leaders

FEBRUARY 2026

What We'll Cover

Six chapters. One question.

01

THE AI LANDSCAPE

Slides 3-4

02

HOW AI WORKS

Slides 5-6

03

THE LIMITATIONS

Slides 7-9

04

THE REALITY CHECK

Slides 10-12

05

SUBQ'S APPROACH

Slides 13-16

06

THE QUESTION

Slides 17-18

AI Is Everywhere

AI is no longer a future technology — it's in every newsroom, every workflow, every phone.

900M

weekly active ChatGPT users

80%

of enterprises have adopted some AI tool

42%

of companies abandoned most AI initiatives in 2025

900M	80%	42%
Users	Adoption	Abandoned

Media Is Ground Zero

Your industry's core outputs — text, images, audio, video — are exactly what these models natively process.

ALREADY HAPPENING

The AP has used automated earnings reporting since 2014 — 3,000+ stories per quarter, written by AI.

Reuters Institute 2025: **87%** of newsroom leaders say AI is transforming operations.

THE TENSION

NYT is suing OpenAI over training data. Axel Springer signed a licensing deal. The industry is split between litigation and partnership — often both at once.

THE CORE QUESTION

→ AI needs media content to work.

← Media needs AI to compete.

Neither side has figured out the terms. That's the landscape we're navigating.

Aa

Shared Vocabulary

Four concepts that underpin every AI discussion.

Aa

Next-Token Prediction

Every AI model predicts the most likely next word. Like autocomplete — trained on the entire written internet.

Key: They learn patterns, not facts.

Attention

The 2017 breakthrough powering GPT, Claude, and Gemini. Every word "looks at" every other word simultaneously.

Powerful — but it's why AI gets expensive fast.

Context Window

How much text the AI can "see" at once — its working memory. Ranges from ~3 pages (2020) to ~1,500 pages (2025).

Can it read your full archive? Not yet. Not well.

Hallucination

Confident fabrication of facts. AI predicts plausible-sounding text, not verified truth. It cannot fact-check itself.

CNET, Sports Illustrated learned the hard way.

Impressive — And Dangerous

AI can draft, summarize, and translate at superhuman speed — but cannot reliably fact-check itself.

STRENGTHS

- Fluent writing in any style or tone
- Translation across 100+ languages
- Summarization of long documents
- Pattern recognition across large datasets

WEAKNESSES

- Hallucination — confident fabrication of facts
- No memory between conversations
- Cannot verify its own claims
- Training data has a cutoff — not real-time

AP: 3,000+ automated earnings stories/quarter

CNET: AI articles riddled with errors

Sports Illustrated: AI articles under fake bios

The AI's Working Memory

Context windows — and why bigger doesn't mean better

The context window has grown 500× in five years:

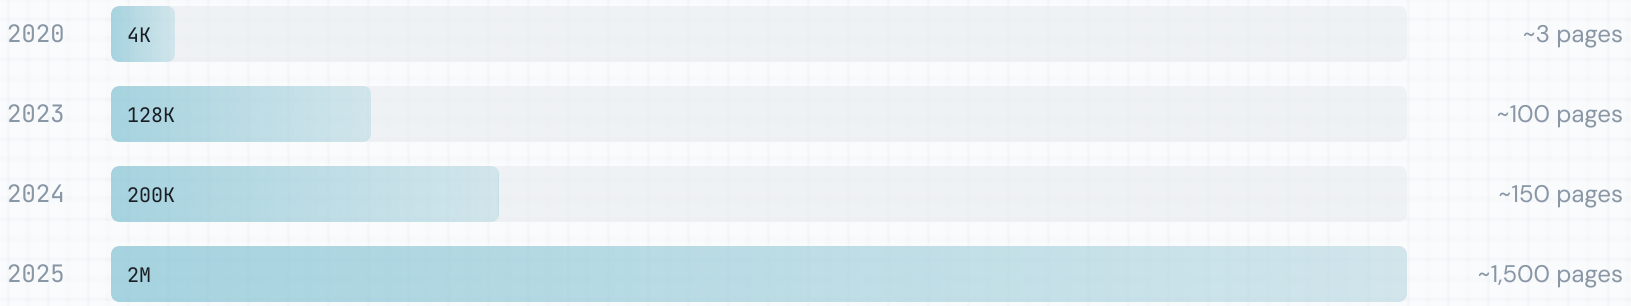

The catch: Like reading 1,500 pages and answering about page 743. Stanford found accuracy dropped to **56%** for information in the middle — worse than not reading at all.

FOR MEDIA

- Can it read an entire investigative file?
- A full day's wire copy?
- Your complete archive?

→ Not yet. Not well.

Why AI Gets Expensive Fast

The cost wall that limits every AI system today

Remember “attention”? Every word must look at every other word. The cost grows as the **square** of the document length.

2× the text
= 4× the cost

10× the text
= 100× the cost

This is not a software bug — it's the fundamental math. And it's **the bottleneck** limiting every AI system you use today.

THE NEWSROOM ANALOGY

10 reporters × 10
= 100 briefings ✓

100 reporters × 100
= 10,000 briefings &x26A0;

1,000 reporters × 1,000
= 1,000,000 briefings &x2717;

Every reporter briefing every other. Works at 10. Breaks at 1,000.

Search-Then-Read

How Perplexity and most AI research tools work today

Question — User query → **Search** — Your archive → **Retrieve** — Top snippets → **Generate** — AI reads snippets → **Answer** — With citations

WHY MEDIA USES IT

- Perplexity: $200M ARR, 780M monthly queries
- Can process archives much larger than any context window
- Provides source citations (critical for journalism)

LIMITATIONS

- Only as good as the search step — miss the right doc, miss the answer
- Misses connections across documents
- Treats content as isolated chunks, not a connected whole
- 50+ copyright lawsuits against Perplexity alone

95% of AI Deployments Show No ROI

Challapally, Pease, Raskar, Chari — "The GenAI Divide" — MIT NANDA, July 2025

52 executive interviews. 153 surveys. 300+ implementations reviewed.

95% of enterprise AI pilots deliver zero measurable business return.

THE DEPLOYMENT FUNNEL

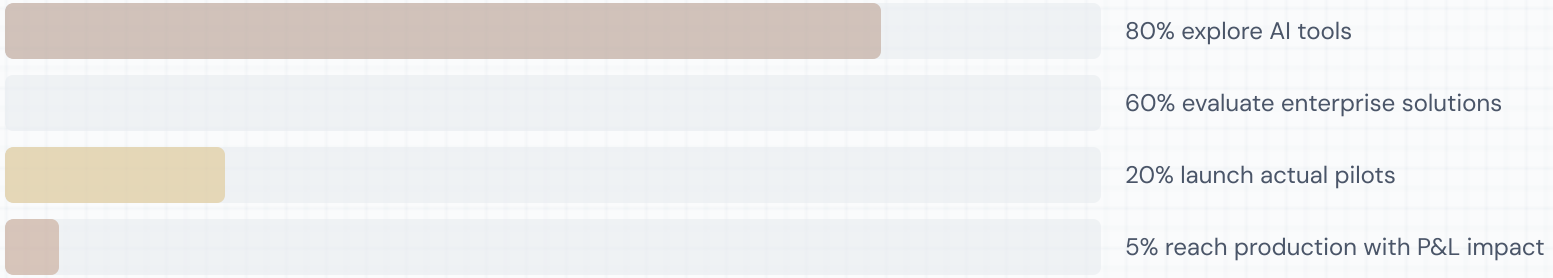

CORROBORATED BY

- S&P Global 2025: 42% abandoned most AI initiatives
- RAND Corp: 80% AI project failure rate — 2× traditional IT

NOTABLE FINDING

Media & Telecom is **1 of 2** sectors showing real structural transformation. The opportunity is real — the execution gap is massive.

95%	52	5%
No ROI	Executives	Production

Why Deployments Fail

The root cause is not model quality — it's organizational and architectural

Generic tools don't adapt

ChatGPT works for individuals but stalls in enterprise: no memory, no workflow integration, no organizational learning.

Too expensive to scale

The cost wall means processing your full archive is prohibitively expensive. Companies settle for snippet tools that miss connections.

Shadow AI economy

90% of employees use personal AI daily, bypassing stalled official initiatives. Transformation is happening — just not through IT.

“It's not the quality of the AI models, but the learning gap for both tools and organizations.”

— Aditya Challapally, MIT NANDA

The Media-Specific Challenge

Deployment failure + unresolved IP and data sovereignty questions.

IN LITIGATION

NYT vs. OpenAI — court ordered 20M ChatGPT logs (Jan 2026)

The Intercept, Raw Story vs. OpenAI

Getty Images vs. Stability AI

TRAFFIC IMPACT

Google AI Overviews: **61% decline** in organic click-through for publishers whose content appears in AI-generated answers.

LICENSING DEALS

News Corp × OpenAI — $250M over 5 years

Associated Press, Axel Springer × OpenAI

Vox Media, The Atlantic × OpenAI

THE DATA QUESTION

When you use an AI tool, does your content train their model? Can you control where your data goes?

For media with sensitive sources and proprietary archives — data sovereignty isn't optional.

What If the Cost Wall Is the Root Cause?

01 Deployments fail because tools can't process enough context

Search-and-snippet tools miss cross-document connections. Full-document processing is prohibitively expensive.

02 They can't process enough context because attention is too expensive

The cost wall means enterprises settle for fragmented, snippet-based tools.

03 Break the cost wall → unlock robust deployment

Make long-document processing 10-40× cheaper, and deployment becomes viable infrastructure.

The 5% that succeed share one thing: tools that can actually learn from — and operate on — the organization's full content.

SubQ

A Different Way to Process Information

SubQ isn't a wrapper around ChatGPT. It rethinks how AI reads documents — breaking through the cost wall.

PROBLEM

Every word checking every other word — cost explodes

INSIGHT

Not every section needs to attend to every other — focus where it matters

RESULT

5–8× faster on long documents at a fraction of the cost

How SubQ Works

The newsroom analogy, continued

STANDARD AI

Every reporter briefs every other reporter.

10 reporters = 100 briefings
100 reporters = 10,000 briefings
1,000 reporters = 1,000,000 briefings

SUBQ

Each reporter only briefs those on related stories.

10 reporters = ~80 briefings
100 reporters = ~800 briefings
1,000 reporters = ~8,000 briefings

- Full document collections
- 5–8× faster (benchmarked)
- Cost grows slowly
- Your data never leaves

SubQ vs. Standard Approaches

	SEARCH-THEN-READ	BRUTE-FORCE READING	SUBQ
How it reads	Searches, then reads snippets	Reads everything at once	Reads everything, focuses on what matters
Cost at scale	Low (search only)	Explodes with length	Grows slowly
Cross-document connections	—	✓	✓
Full document comprehension	—	Drops to 56% mid-doc	✓
Source citations	✓	—	✓
Data sovereignty	Depends on vendor	Cloud-only	Your infrastructure

Not “smarter AI” — **smarter architecture**. The same intelligence, applied where it matters.

Where the Opportunities Are

Three areas where AI that reads everything — not just snippets — changes what's possible.

Investigative Journalism

AI that cross-references thousands of documents and finds patterns humans miss. Think: Panama Papers at 10× speed — with every source cited.

Archive Monetization

Decades of content made searchable and contextually accessible. True comprehension of your full library, on your infrastructure.

Fact-Checking at Scale

Verify claims against your own published record in real-time. Your archive becomes your fact-checking engine.

The "answer engine" market: $2.3B today → $7.3B by 2030. The question is who controls the infrastructure.

INVESTIGATE

SubQ

Where can AI that reads everything — not just snippets — change how you work?

SCALE

ARCHIVE

The 5% of successful deployments start with the workflow, not the technology.

You know the workflows. We built the architecture.

What are the highest-friction points in your current AI workflows?

We'd like to explore that — not as a pitch, but as a diagnostic.

PERSONALIZE

VERIFY